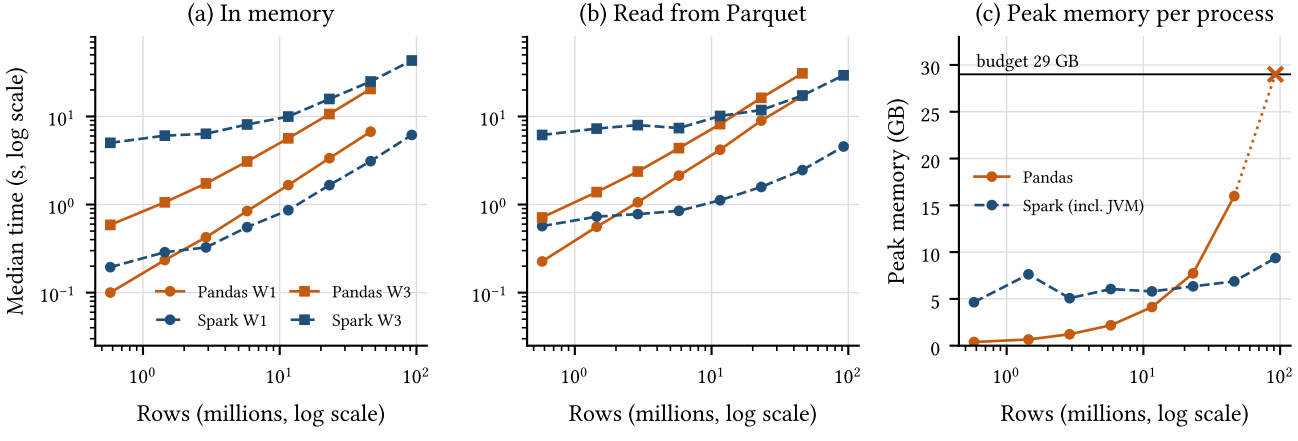
(a) In memory
(b) Read from Parquet
(c) Peak memory per process
Median time (s, log scale)
Peak memory (GB)
Rows (millions, log scale)
Rows (millions, log scale)
Rows (millions, log scale)
Pandas W1
Pandas W3
Spark W1
Spark W3
budget 29 GB
Pandas
Spark (incl. JVM)