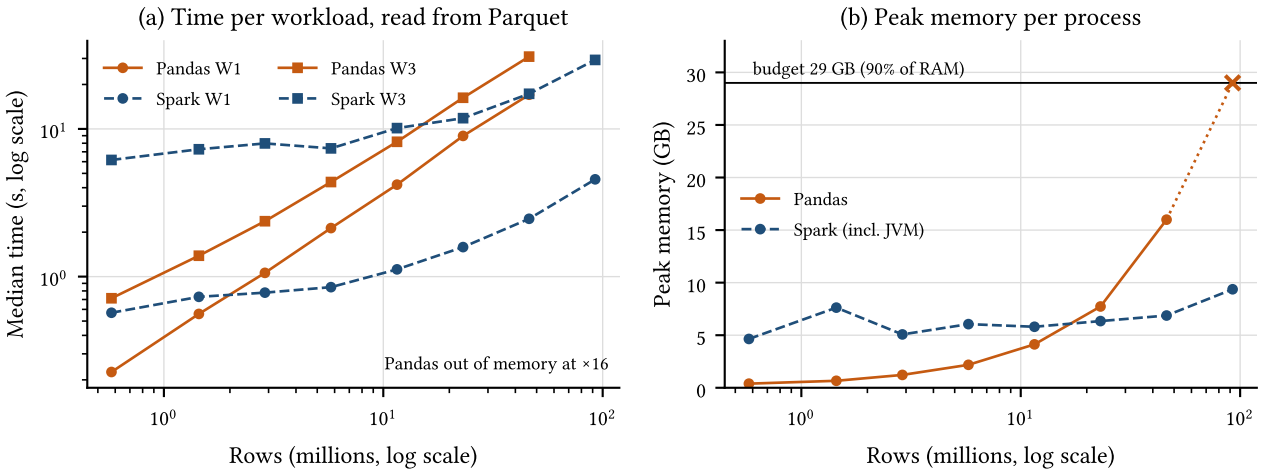
(a) Time per workload, read from Parquet
Pandas W1
Pandas W3
Spark W1
Spark W3
Pandas out of memory at ×16
Median time (s, log scale)
Rows (millions, log scale)
(b) Peak memory per process
budget 29 GB (90% of RAM)
Pandas
Spark (incl. JVM)
Peak memory (GB)
Rows (millions, log scale)